

# R Workshop

April 15, 2021

## 1 R Coding Workshop

4/15/2021

### 1.1 Overview

Graduate Student Instructor: Kayleigh Barnes

Email: [kayleighnb@berkeley.edu](mailto:kayleighnb@berkeley.edu)

#### 1.1.1 Goals for today

This session is intended to guide you through the practical implementation of basic analytic techniques in R in Jupyter notebooks. R is an open-source statistical computing software used to analyze data. A Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and narrative text. This workshop will be focused on interactive demonstration in R, but also include time for additional questions and guidance in working through the sample code. We will cover some fundamental coding techniques that will help you in Econ 140, basic data science classes, or research assistant positions. This workshop is for *beginners* that have little or no coding experience.

#### 1.1.2 Important notes

- One attendee from today's workshop will be randomly selected to win a 20 dollar gift card to Amazon
- Attendance to this workshop comes with free access to datacamp through July. Datacamp offers online courses in both R and Python so that you can continue learning after today's workshop
- Link to join Berkeley Econ's datacamp group with @berkeley.edu ID: [here](#) (make sure you're signed out of datacamp before clicking this - otherwise the sign-up gets screwed and you'll be asked to pay after the first chapter of any course)

### 1.2 Jupyter and R Basics

- To create a new notebook, click the "New" button and select R

- Write R script by selecting the option “Code” from the dropdown list, or write text by selecting “Markdown”
- Select “Insert” to add a block of text or code
- Run code by highlighting and selecting “Run”
- Use the # symbol to add comments to the script, or to add headlines to text selections
- To clear your coding output, select Cell=>All Output=>Clear

```
[1]: # Clear the workspace, this removes all data and numbers you have stored or
      ↪ saved in R
      rm(list = ls())

      # The help function, using ? or help() before a command will bring up
      ↪ information on what the command does
      ?setwd
      help(setwd)
```

```
[2]: #The working directory is the location that R will look for data in
      # this is the same as telling your computer to look in a documents folder
      ↪ when uploading soemthing
      getwd()
```

```
'/home/jovyan/my-work'
```

User written open-source packages are needed for specific functionality in R (e.g. nice graphics). However, we need to manually install these packages (once) and load them at the beginning of every script. Packages have been pre-installed in Jupyter notebooks. If you are wondering why a command you’ve used before is no longer working, it may be because you haven’t loaded the package.

```
[3]: #Install packages
      install.packages('ggplot2')

      # Load required packages
      library(ggplot2)
```

```
Updating HTML index of packages in '.Library'
```

```
Making 'packages.html' ...
done
```

### 1.3 Loading in data and summary statistics

Now let’s load in the data set. Make sure you have uploaded the data to Jupyter before running the next line of code. We are going to use data on a set of households in Mexico in the 1990’s. The data includes a village ID, a household ID, and demogrpahic variables like income, household size, age and gender of the head of household and a poverty indicator.

```
[4]: # Reading data into R from a CSV file
# ?read.table # delete the # at the beginning of this line to view the help
# entry for the "read" command
MyFirstData <- read.csv('MyFirstData.csv', header = TRUE)
```

Notice that there is no output from the code that reads in the data. Unlike excel, R stores the data in the background and we need to use specific commands to interact with it. Once it's read in, we can use several commands to describe the data

```
[5]: # Structure of the Data
str(MyFirstData)
```

```
'data.frame': 1200 obs. of 8 variables:
 $ villid : int 1001106 1001106 1001106 1001106 1001106 1001106 1001106 1001106
1001106 1001106 1001106 ...
 $ hogid : Factor w/ 1200 levels "0101103002.0639",...: 10 11 12 13 14 15 16
17 18 19 ...
 $ D_HH : int 1 1 1 1 1 1 1 1 1 1 ...
 $ IncomeLab: int NA NA NA 3200 NA 4320 4800 NA NA 3200 ...
 $ famsize : int 6 6 6 5 5 5 5 6 6 3 ...
 $ agehead : int 29 43 43 25 40 40 39 45 42 22 ...
 $ sexhead : Factor w/ 2 levels "Female","Male": 1 2 2 1 2 2 2 1 2 2 ...
 $ pov_HH : Factor w/ 2 levels "no pobre","pobre": 2 1 1 2 1 1 1 2 1 2 ...
```

```
[6]: # Summary of the Data
summary(MyFirstData)
```

villid	hogid	D_HH	IncomeLab
Min. :1001106	0101103002.0639:	1	Min. :0.0000
1st Qu.:7011004	0101103004.0640:	1	1st Qu.:1.0000
Median :7011019	0101103006.0641:	1	Median :1.0000
Mean :5951112	0101103008.0642:	1	Mean :0.8107
3rd Qu.:7015003	0101103012.0644:	1	3rd Qu.:1.0000
Max. :7015038	0101103014.0645:	1	Max. :1.0000
	(Other) :1194	NA's :17	NA's :1024
famsize	agehead	sexhead	pov_HH
Min. : 1.000	Min. :16.00	Female:257	no pobre:203
1st Qu.: 3.000	1st Qu.:37.00	Male :943	pobre :996
Median : 5.000	Median :47.00		NA's : 1
Mean : 4.801	Mean :49.12		
3rd Qu.: 6.000	3rd Qu.:60.00		
Max. :16.000	Max. :96.00		
	NA's :1		

```
[7]: # Variable Names
colnames(MyFirstData)
```

1. 'villid' 2. 'hogid' 3. 'D\_HH' 4. 'IncomeLab' 5. 'famsize' 6. 'agehead' 7. 'sexhead' 8. 'pov\_HH'

```
[8]: #Number of Observations
      nrow(MyFirstData)
```

1200

```
[9]: #Display first 6 rows of the data
      head(MyFirstData)
```

	villid	hogid	D_HH	IncomeLab	famsize	agehead	sexhead	pov_HH
	<int>	<fct>	<int>	<int>	<int>	<int>	<fct>	<fct>
1	1001106	0101103050.0539	1	NA	6	29	Female	pobre
2	1001106	0101103052.0540	1	NA	6	43	Male	no pobre
3	1001106	0101103054.0541	1	NA	6	43	Male	no pobre
4	1001106	0101103056.0542	1	3200	5	25	Female	pobre
5	1001106	0101103058.0543	1	NA	5	40	Male	no pobre
6	1001106	0101103060.0544	1	4320	5	40	Male	no pobre

A data.frame: 6 × 8

```
[10]: #Tabulate a specific variable (to refer to a variable, use Dataset$VariableName)
       table(MyFirstData$sexhead)
```

Female	Male
257	943

## 1.4 Basic Data Cleaning and Formatting

### 1.4.1 Category Variable

Right now, we have two categorical variables: sexhead, which indicates the sex of the head of household and pov\_HH, which indicates whether a household is below the poverty line. The data entries for these variables are text rather than numbers (we call these string variables in the data science world). Often when doing data analysis, it is easier to map categorical text variables to numbers, particularly 0 and 1. These variables that contain only 0's and 1's are called dummy variables.

Now, suppose we want to create a poor\_male variable, which will be defined as 1 if the household is categorized as poor (pov\_HH = pobre) and the head of the household is male (sexhead is Male), and 0 otherwise.

```
[11]: #Create one dummy variable based on T/F condition
MyFirstData$poor_male <- ifelse(MyFirstData$pov_HH == 'pobre' &
  ↪MyFirstData$sexhead == 'Male', 1, 0)
#tabulate the observations
table(MyFirstData$poor_male)
```

0	1
413	786

### 1.4.2 Numerical Variable

We can use regular mathematical operations to create numerical variables from other variables.

```
[12]: #Squaring an existing variable
MyFirstData$agehead2 <- MyFirstData$agehead^2
summary(MyFirstData$agehead2)

#Creating a constant
MyFirstData$constant <- 1
summary(MyFirstData$constant)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
256	1369	2209	2656	3600	9216	1

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1	1	1	1	1	1

### New Datasets We may also want to create a new data that summarizes the old, or is a subset of the original dataset.

```
[13]: #Subset of only observations with male head of hh
data_males<-MyFirstData[ which(MyFirstData$sexhead=='Male'),]
summary(data_males)

#First select variables to aggregate
myvars <- c("villid", "IncomeLab", "famsize", "agehead")
meandata <- MyFirstData[myvars]

#Collapse data to get average values by village. Could also use "sum" as the
→function to get totals
meandata<-aggregate(meandata, by = list(meandata$villid), FUN = mean)
nrow(meandata)
summary(meandata)
```

villid		hogid		D_HH		IncomeLab	
Min.	:1001106	0101103002.0639:	1	Min.	:0.0000	Min.	: 160
1st Qu.:	7011004	0101103004.0640:	1	1st Qu.:	1.0000	1st Qu.:	1388
Median	:7011019	0101103006.0641:	1	Median	:1.0000	Median	: 1600
Mean	:5954876	0101103008.0642:	1	Mean	:0.8254	Mean	: 2277
3rd Qu.:	7015003	0101103012.0644:	1	3rd Qu.:	1.0000	3rd Qu.:	2850
Max.	:7015038	0101103018.0647:	1	Max.	:1.0000	Max.	:27000
		(Other)	:937	NA's	:15	NA's	:807
famsize		agehead		sexhead		pov_HH	
						poor_male	

Min. : 1.000	Min. :18.00	Female: 0	no pobre:156	Min. :0.0000
1st Qu.: 4.000	1st Qu.:36.25	Male :943	pobre :786	1st Qu.:1.0000
Median : 5.000	Median :47.00		NA's : 1	Median :1.0000
Mean : 5.022	Mean :49.05			Mean :0.8344
3rd Qu.: 6.000	3rd Qu.:60.00			3rd Qu.:1.0000
Max. :16.000	Max. :94.00			Max. :1.0000
	NA's :1			NA's :1

agehead2	constant
Min. : 324	Min. :1
1st Qu.:1314	1st Qu.:1
Median :2209	Median :1
Mean :2638	Mean :1
3rd Qu.:3600	3rd Qu.:1
Max. :8836	Max. :1
NA's :1	

24

Group.1	villid	IncomeLab	famsize
Min. :1001106	Min. :1001106	Min. : NA	Min. :4.504
1st Qu.:1002032	1st Qu.:1002032	1st Qu.: NA	1st Qu.:4.721
Median :1008514	Median :1008514	Median : NA	Median :5.348
Mean :2506884	Mean :2506884	Mean :NaN	Mean :5.316
3rd Qu.:2509100	3rd Qu.:2509100	3rd Qu.: NA	3rd Qu.:5.688
Max. :7015038	Max. :7015038	Max. : NA	Max. :6.800
		NA's :24	

agehead
Min. :36.00
1st Qu.:40.96
Median :45.58
Mean :45.34
3rd Qu.:50.25
Max. :52.27
NA's :1

## 1.5 Making comparisons - T-Tests

A main goal of working with data is to make inferences about the population we are interested in. Much of Econ 140 will be focused on methods to make these inferences: What is the relationship between two variables? Did an experiment have a significant treatment effect?

If you have taken Stats 20, you are likely already familiar with a t-test. T-tests compare the difference in the means of a variable between two groups. The test statistic tells us whether the difference is *significant*, that is we can confidently say that the two groups are different.

```
[14]: #let's run a t-test comparing the average family size for households above and  
↪below the poverty line  
t.test(MyFirstData$famsize ~ MyFirstData$pov_HH, var.equal = TRUE)
```

Two Sample t-test

```
data: MyFirstData$famsize by MyFirstData$pov_HH  
t = -5.2032, df = 1197, p-value = 2.303e-07  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 -1.221131 -0.552397  
sample estimates:  
mean in group no pobre    mean in group pobre  
      4.064039             4.950803
```


## 1.6 Visualizing Data

Make sure that the ggplot2 package is included at the top of the script. Below, we show an example of a scatterplot using ggplot. “geom” can be used to denote different types of graphs such as a line graph.

```
[15]: ggplot(MyFirstData, aes(x = agehead, y=famsize)) + geom_point()  
      ?geom_line
```

Warning message:

"Removed 1 rows containing missing values (geom\_point)."




We can use a direct function or ggplot to create a histogram. Notice that changing the options in the function allows you to customize the graph. Use the help function to learn more about the options for each command.

```
[16]: # Base Graphics
hist(MyFirstData$agehead)
hist(MyFirstData$agehead, col = "blue", main = "Histogram of age")
# ggplot2
ggplot(MyFirstData, aes(x = agehead)) + geom_histogram(fill = "blue") +
  ggtitle("Histogram of age")
```



**Histogram of MyFirstData\$agehead**




``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.

Warning message:

"Removed 1 rows containing non-finite values (stat\_bin)."





[ ]: